

Measuring Student Learning in the Age of AI: A Framework and Recommendations for BYU-Idaho

Prepared by: Mat Miles · **Requested by:** the BYU Provo Library Committee · **Institution:** Brigham Young University-Idaho · **Date:** July 2026

Status: Full draft (v8, terse, no dashes; reframed for BYU-Idaho, adds current policy and website recommendations). Citation style: APA 7th ed.

References verified. Evidence base: `02\_Research\_Notes\_and\_Sources.md`.

Executive Summary

Generative AI changed the economics of assessment. A credible essay, problem set, or block of code now takes seconds to produce, so the student product no longer proves what a student knows. The common reflex, detect and ban, has failed. AI detectors are inaccurate, biased against non-native English, Black, and neurodivergent students, easily evaded, and unable to prove their own results (Liang et al., 2023; Weber-Wulff et al., 2023; Ardito, 2025). Regulators and major universities, including US institutions such as Vanderbilt, are dropping them.

The field has reframed the problem from integrity to validity: design assessment that measures learning rather than police a tool that is now everywhere (Corbin et al., 2026; Lodge et al., 2023). This report synthesizes 2023 to 2026 scholarship and recommends a two-horizon path. It is written for BYU-Idaho and measures the university's current policy against the evidence (Section 7).

Short-term wins, now: stop using detectors as evidence of misconduct; add a few secured (in-person or oral) checkpoints that verify individual understanding; state permitted AI use for each task with a shared AI-use scale; weight the process, not just the product.

Long-term adjustments: move assessment security to the program level, securing a few key moments rather than every task, for example a secured oral at the end of each year plus a proctored capstone; redesign in teams, not alone, for example a department building one shared assessment with a librarian and an instructional designer; sustain a shared faculty resource repository, for example a department site holding vetted task templates, rubrics, and sample AI-use statements.

Done well, this serves three aims at once: valid assessment of learning; stronger teacher-student relationships, since oral formats build understanding and trust together; and graduates ready to use AI well in work, community, and home. The report ends in a best-practices repository: curated practices, leading figures, and ready-to-use resources, plus worked examples (Appendices A to G). BYU-Idaho is well placed: it already avoids AI detectors and requires disclosure, so this report extends a sound base rather than reversing course. We ask the

committee and administration to endorse the framework, resource the repository, and authorize a program-level pilot.

1. Introduction & Problem Statement

The project's purpose: find ways to assess knowledge and skills across BYU-Idaho now that generative AI has reshaped the work. The shift is economic. Producing a plausible artifact, whether an essay, a lab report, or working code, is fast and nearly free, so these products no longer measure many kinds of learning (Desai, 2025). Use is near-universal. Of 8,000-plus Australian students, 83% reported using AI, 40% when they knew they should not; a 2025 UK survey of 1,041 students found 88% used AI in assessment (Freeman, 2025). The US picture is similar. Faculty use is broad but still largely exploratory (Baytas & Ruediger, 2025), and institutions are worried and underprepared: in the EDUCAUSE 2025 landscape study (n = 788), 55% expect more academic dishonesty and 56% more student over-reliance by 2027, yet only 22% have an institution-wide AI approach (EDUCAUSE, 2025).

AI did not create the assessment problem. It exposed one. No university ever had "perfect insight into what students know... That's always been a myth" (Corbin et al., 2026). The product was always an imperfect proxy; cheap generation made the gap undeniable. The question shifts from "how do we stop AI cheating?" to "what is worth measuring, and how do we measure it credibly?"

A second problem compounds the first. Faculty and students are unsure where legitimate AI use ends and cheating begins. Both "draw boundaries regarding acceptable AI use which they feel they cannot defend with confidence," and once assessment is not trusted, neither are the degrees it certifies (Corbin et al., 2025). Clarifying that line is part of the work, not a precondition.

This report treats deskilling versus upskilling as a design problem, not a ban-versus-allow fight. It pursues two horizons, short-term wins and long-term adjustments, toward three aims: valid assessment of learning; stronger teacher-student relationships; graduates ready to serve in work, community, and home. Faculty are the leverage point, since assessment is built course by course. Sections follow: evidence (2), trends (3), framework (4), the repository (5), risks (6), where BYU-Idaho stands today (7), recommendations (8).

2. Literature Review

2.1 How generative AI disrupts assessment validity

AI is a validity problem, not just an integrity one. When AI can generate the artifact, a grade "may increasingly reflect a combination of individual understanding and external cognitive

support rather than solely independent competence" (Brattli et al., 2026). A natural experiment shows it. Across five years of one compulsory course (N = 1,066), the shift from an AI-permissive take-home exam (2021 to 2024) to an AI-restricted format (2025) sharply raised failure rates and cut mid-range grades while top grades held: $\chi^2(5, N = 1066) = 60.62, p < .001$, Cramér's V = 0.24. The authors conclude the two formats "may not be measurement-equivalent under conditions of widespread GenAI use," threatening construct validity (Brattli et al., 2026). One course limits generalizability, but it grounds an otherwise conceptual claim.

The field calls AI-and-assessment a "wicked problem" meeting all ten of Rittel and Webber's (1973) criteria: no correct solution, only better or worse responses, best navigated rather than solved (Corbin et al., 2026).

2.2 Academic integrity in the AI era: beyond detection and banning

Detection as enforcement is the literature's strongest, most consistent failure.

- **Bias.** Seven GPT detectors misflagged 61.3% of non-native (TOEFL) essays as AI while scoring native essays near-perfectly; one flagged 97.8%, and all unanimously misflagged 19.8%. The authors "strongly caution against the use of GPT detectors in evaluative or educational settings" (Liang et al., 2023). Black and neurodivergent students face similar harm (Deep et al., 2025).
- **Unreliability.** A test of 14 tools found them "neither accurate nor reliable," biased toward calling text human (Weber-Wulff et al., 2023).
- **Evasion.** Paraphrasing, style prompts, and parameter tweaks defeat detectors; new bypasses spread within days (Ardito, 2025).
- **False accusation at scale.** At a 1% false-positive rate, an honest student writing 70 essays has a 50.5% chance of being wrongly flagged at least once (Ardito, 2025).
- **Non-falsifiability.** A detector "provides no evidence other than the results of its assessment," and since an LLM can output any text, even the prompts cannot prove authorship (Ardito, 2025).
- **Incompatible with teaching AI.** Detection cannot tell good use from bad, so it cannot coexist with permitting or teaching AI (Ardito, 2025).

Scholarship has moved from containment (2023) to a "pedagogical-trust model" favoring ethical engagement over surveillance (Slimi, 2026). UCLA, University of California campuses, Manchester, and Vanderbilt have dropped or discouraged detectors; TEQSA calls reliable detection "all but impossible" (Lodge et al., 2023; TEQSA, 2025). In the US this moved from opinion to policy early: Vanderbilt disabled Turnitin's AI detector in 2023, citing false positives and bias against non-native English writers (Vanderbilt University, 2023), and the University of Pittsburgh's teaching center did the same, judging detection not reliable enough to deploy (University of Pittsburgh, 2023).

The constructive alternative to detection is setting clear, task-level expectations for AI use. The most adopted tool is the AI Assessment Scale (AIAS), which sets the allowed level of AI for each task across five levels, from No AI to AI Exploration (M. Perkins et al., 2024); similar four-level

scales exist (Stokkink, 2025). The AIAS has spread to hundreds of institutions and 30-plus languages and is cited by TEQSA, and a single-site pilot reported higher attainment and fewer misconduct cases (Furze et al., 2024). For faculty, the practical gain is clarity: stating the allowed AI level for each task removes ambiguity and gives students and markers one shared standard, which is what makes the pilot's drop in misconduct plausible. Read this as promising-but-early: the framework is widely used, but strong impact evidence rests on one pilot, and declarations alone can fall flat. In Scottish English for Academic Purposes (EAP) departments, cover-sheet AI declarations were seen as an ineffective stop-gap, while context-bound tasks and named marking were felt to reduce misconduct (Urquhart & Ngo, 2026). The AIAS authors dropped the old red-to-green "traffic light" colours in 2024, so it is no longer a stoplight tool.

2.3 AI-resistant / AI-robust assessment design

The leading design response separates secured from open assessment. Sydney's "two-lane" framework makes Lane 1 secured, in-person assessment of learning, mostly at program level for accreditation, and Lane 2 open, unsecured assessment for and as learning, mostly at unit level, assuming AI use; from January 2025 Sydney coordinators cannot prohibit AI in open tasks (University of Sydney, 2024). The logic: "any unenforceable restriction... damages assessment validity," so integrity comes from securing Lane 1, not policing Lane 2 (University of Sydney, 2024). TEQSA treats assessment security as programmatic: secure key moments across a program, not every task (TEQSA, 2025).

Example, a first-year writing course. Lane 1 is a 20-minute in-class handwritten analysis of an unseen passage, which confirms the student can read and reason without help. Lane 2 is a take-home essay where students may use AI to draft, but must submit their prompts and a short paragraph defending their edits. Lane 1 certifies the core skill; Lane 2 builds and assesses skilled AI use.

US teaching centers are adapting this model. The University of Massachusetts Amherst pairs "AI-immune" tasks (in-class writing, reflection, presentations) with "AI-integrated" tasks whose outcomes explicitly involve AI, citing the Sydney framework (University of Massachusetts Amherst, n.d.). Note that "AI-immune" is aspirational, since nothing is fully AI-proof; the secured/open logic is the durable idea.

2.4 Authentic, performance, and portfolio assessment

A second strand applies Wiggins's authentic assessment to define "resilient assessment": tasks that mirror real-world challenges and demand understanding AI cannot easily fake, shifting weight from the artifact to the process (G. Perkins, 2026). Formats include portfolios, live projects, and pitch-style presentations grounded in personal experience that resist outsourcing. This links validity (aim 1) to workplace capability (aim 3).

Performance examples: a nursing student runs a simulated patient handover; a marketing student pitches a live campaign to a panel; an education student teaches a 10-minute lesson and is graded on the teaching, not a written plan. Portfolio examples: an engineering student keeps a semester portfolio of design iterations, each with a short rationale; a writing student

submits a folder of drafts that shows revision over time; an art student presents a body of work plus a reflective statement on each choice.

One caution and two extensions. Caution: authentic tasks are not automatically AI-proof. Across two UK universities, markers generally could not tell AI-assisted work from unassisted, and more authentic tasks were no easier to spot (Kofinas et al., 2025), so authenticity helps validity but does not by itself secure a task. Extensions, still early: self-assessment and reflective commentary, where students judge and explain their own work; and purpose-driven and role-play tasks, including role-play with an AI persona, for example a student advising a simulated AI client. Adoption is rising but rigorous impact evidence is thin (Weng et al., 2024).

US practice is converging on the same ground. AAC&U's VALUE rubrics give a shared, outcomes-based way to judge authentic work (Association of American Colleges and Universities, n.d.), and the widely used guide Teaching with AI argues for redesigning around human learning outcomes (Bowen & Watson, 2024). Carnegie Mellon's GAITAR initiative pushes backward design from outcomes but treats whether AI helps or harms learning as an open, measurable question, not a settled one (Carnegie Mellon University, 2023).

2.5 Oral, viva, and interactive assessment

Oral and interactive oral assessments (IOAs) make students explain and defend their understanding live, resisting AI substitution while building professional skill (G. Perkins, 2026). They pay a relational dividend: a 15-minute conversation can reveal understanding better than a written report and give faster, more personal feedback (University of Melbourne, n.d.). That serves aim 2, recasting oral assessment as relationship-building, not just verification. Its cost is workload (see 2.7 and 6).

Examples: a 10-minute viva where a student explains and defends their submitted lab report; a 5-minute "explain your code" check at the terminal, where the student walks through a function and answers one change request; an interactive oral where the instructor poses an unscripted workplace scenario and the student responds in role; a "two questions at submission" check, where each student answers two quick questions about their own essay when handing it in.

Adoption is real: one Sydney STEM unit ran 400-plus interactive orals, each a structured 10-minute conversation, as a live alternative to AI-compromised written work (University of Sydney, 2025). Treat this as an adoption signal, not proven validity; there is not yet evidence that the oral scores tracked performance on the secured exam. Emerging AI tools aim to cut the workload cost. Voice agents let students rehearse verbal scenarios for formative practice. AI assessment agents, or AI-assisted vivas, watch a student work and then ask follow-up questions to have them explain what they did and why. Early pilots exist, such as a Socratic AI oral tool, but evidence is thin and mostly vendor-claimed or single-site (see Appendix F). US work shows both the promise and the limits: a Georgia Tech platform scaled AI-run Socratic oral checks to thousands of students (Hung et al., 2024), with later developer-reported learning gains that are not yet independently replicated; and a separate proof-of-concept runs an LLM viva on a

student's own submission (Church et al., 2025). The first is promising-but-early, the second unproven. Use these tools to practice and prepare, not yet as the sole basis for a grade.

2.6 Competency-based and programmatic assessment

No single task can certify learning under widespread AI, so scholars anchor capability judgments in programmatic portfolios or longitudinal records rather than one task (Corbin et al., 2025). TEQSA likewise wants "multiple, inclusive and contextualised" evidence and spreads secured checkpoints across a program instead of one capstone (TEQSA, 2025; University of Sydney, 2024). The one-shot summative exam becomes one input, not the sole guarantee.

Examples: a program maps one core competency, say evidence-based argument, across three courses and certifies it with a rising-difficulty portfolio rather than one final essay; a degree adds a short secured oral at years one, two, and three to confirm progression; a clinical program certifies skills through an end-of-year OSCE (objective structured clinical examination) instead of per-unit essays; a major replaces a single capstone exam with a defended portfolio plus one proctored task.

Two program-level shifts follow. First, rewrite course and program outcomes for an AI world, so they name the judgment and process students must show, not just a finished product. A scoping review of 34 studies concludes traditional methods no longer work well and sorts the responses into refocused and AI-incorporated assessment (Weng et al., 2024). Second, move from a few high-stakes exams toward frequent low-stakes formative work, assessment for learning rather than only of learning. Students who used AI to build and augment knowledge, a mastery approach, reached higher-order outcomes than those who used it mechanically (Pallant et al., 2025). Both are promising-but-early, and practice lags the advice: many departments still tweak old tasks rather than redesign them (Urquhart & Ngo, 2026).

2.7 Sustainable assessment redesign

Workload is the binding constraint. Oral formats are costly: 250 students at 10 minutes each is 2,500 minutes of staff time. Nicola-Richmond et al. (2025) argue redesign needs program-wide teams pooling disciplinary, AI, and assessment expertise, since no one person holds it all; this makes redesign an institutional investment, not an individual burden. Laidlaw (2026) adds that redesign must stay equitable and respect staff capacity.

How to implement: form a small redesign team per program (a coordinator, a discipline expert, an instructional designer); start with one high-enrollment course and copy what works; reuse a shared bank of secured-task templates rather than building each from scratch; keep orals short, 5 to 10 minutes on a sample of work, instead of full vivas for everyone; stagger oral checkpoints across weeks to spread the load; attach orals to existing tutorials or labs to avoid new scheduling; and grade with a one-page rubric to keep marking fast and consistent.

2.8 What we are measuring: keeping outcomes meaningful

Two failure modes recur: reverting to exam-heavy formats that "verify learning but fail to promote meaningful retention," and tasks so rigid they test compliance, not capability (TEQSA,

2025). The aim is to protect durable competencies (reasoning, writing, communication, evaluative judgment, ethical reasoning) while teaching students to use AI as a scaffold for higher-order work, not a substitute (TEQSA, 2025).

3. Current Trends in Assessment

Six trends recur across 2023 to 2026 sources.

1. **From detection to redesign.** Regulators included, the sector has shifted from detection to structural redesign as the only sustainable response (TEQSA, 2025; Slimi, 2026).
2. **Secured checkpoints return, at program level.** In-person, oral, and interactive formats are resurging, but spread across a program rather than as high-stakes single exams (University of Sydney, 2024; TEQSA, 2025).
3. **Policy shifts to "assume AI."** Sydney's stance typifies the move from prohibition to scaffolded, disclosed use; TEQSA required every Australian provider to file a GenAI action plan in 2024.
4. **Expectation-setting scales spread.** Task-level AI-use scales, led by the five-level AI Assessment Scale, are widely adopted to make the rules explicit (M. Perkins et al., 2024).
5. **The formative shift.** Programs move from a few high-stakes exams toward frequent low-stakes formative work (Weng et al., 2024).
6. **Tensions persist:** equity and access, faculty workload, the validity of "secure" tasks, and tool churn. Govern them openly; do not declare them solved.

In the US, regional accreditors have begun to engage: a 2025 statement from the Council of Regional Accrediting Commissions addresses using AI to advance learning, evaluation, and recognition (Council of Regional Accrediting Commissions, 2025).

4. Conceptual Framework: Measuring Learning When AI Changes the Economics of Assessment

4.1 The core distinction

The frame separates assessing the artifact from assessing the learning, run as secured (AI-controlled) versus open (AI-permitted) tasks. Secured tasks assure individual competence; open tasks build and assess capability with AI, as students will work.

4.2 Competing approaches (evidence-weighted)

Approach	What it assures	Strength of evidence	Key caveat
Two-lane / secure-vs-	Program-level	Strong sector adoption;	Critiqued as too binary

open (Sydney; TEQSA)	assurance plus AI-integrated learning	regulator-backed	(Curtis, 2025)
Fully AI-integrated assessment	Workplace-realistic capability	Conceptually strong; thin efficacy data	Hard to attribute individual competence
Middle-lane (AI permitted with limits)	Graduated, task-specific AI use	Argued as essential (Curtis, 2025)	Boundary-setting burden on faculty
Process / portfolio / ungrading	Development over time	Aligns with validity argument	Workload; consistency of judgment
Oral / authentic	Individual understanding; relationship	Strong practitioner, growing scholarly	Scaling cost; equity and accommodations

A peer-reviewed critique calls a strict all-or-none two-lane model "insupportable": it threatens equity and adds little validity in "secure" tasks, and a middle lane permitting AI with limits is essential (Curtis, 2025). So two-lane is the leading organizing model, not a settled consensus.

4.3 A practical decision method for faculty

For each outcome, faculty ask: must the student show this without AI? If yes, assess it secured (in-person, oral, or interactive). If the outcome is best shown with AI, as most professional work now is, assess it open and grade judgment, process, and critique of AI output. AI stays free for learning and practice either way. This protects a defined AI-free core while teaching and assessing AI-augmented capability elsewhere, grounded in program-level assurance. The two-lane model and the scale work together: two-lane decides which tasks to secure, and the scale names the allowed AI level on each task, with level 1 (No AI) being the secured case. Communicate the result with a shared, task-level AI-use scale, such as the five-level AI Assessment Scale, so students see exactly what is allowed on each task (M. Perkins et al., 2024).

4.4 Designing for validity and workload

Securing every task is unaffordable, so the framework concentrates security at a few high-value checkpoints per program and puts redesign effort into teams (see 2.7). Validity and sustainability are joint constraints, not rivals.

5. Deliverable: Best-Practices Report & Faculty Resource Repository

The deliverable is a living repository that helps faculty improve their own courses.

- **5.1 Curated best practices, by horizon.** Short-term: drop detector-based accusations; add one secured or oral checkpoint per course; state AI rules per task; weight process. Long-

term: program-level assessment mapping; team-based redesign; multiple points of evidence instead of single high-stakes exams.

- **5.2 Leading figures and organizations** to follow (Appendix C).
 - **5.3 Ready-to-use resources** (Appendices B, D, E): redesign templates, process-weighted rubrics, an interactive-oral protocol, an evaluation instrument, and sample AI-policy language with a task-level AI-use scale.
 - **5.4 A self-guided pathway** to map outcomes to secured or open tasks and redesign one assessment.
 - **5.5 Curation and governance** to keep it current as tools and evidence evolve.
 - **5.6 Emerging tools to watch** (Appendix F): voice practice agents and AI-assisted vivas, to pilot, not yet to grade with.
-

6. Implementation Considerations & Risks

- **Equity and access.** Detector bias harms multilingual and minoritized students; secured exams and oral formats need accommodations. Fix: retire detectors; build in accommodations from the start (Laidlaw, 2026).
 - **Workload and buy-in.** Redesign and oral assessment take time. Fix: redesign at program level in teams; secure only a few checkpoints; resource the repository (Nicola-Richmond et al., 2025).
 - **Validity of "secure" tasks.** Over-rigid tasks test compliance (TEQSA, 2025). Fix: prefer authentic secured formats, like oral defense of real work, over rote exams.
 - **Enforcement.** Going beyond detection needs a credible alternative. Fix: secured checkpoints, clear per-task AI rules, an educative not punitive stance.
 - **Tool churn and privacy.** Vendor tools shift and may misuse student data. Fix: assure learning by design, not by tool; review procurement.
 - **Evidence.** Much redesign advocacy rests on small or conceptual studies (Corbin et al., 2025; Slimi, 2026). Fix: pilot with evaluation; treat the program as adaptive.
-

7. Where BYU-Idaho Stands Today: Current Policy and Practice

BYU-Idaho already runs a public AI hub at byui.edu/ai, built on three strategic priorities: advance knowledge and enhance learning, develop workforce readiness, and facilitate student support and advising (BYU-Idaho, n.d.-c). Its approach is faith-centered: AI should support, not replace, human effort, growth, and relationships (BYU-Idaho, n.d.-a, n.d.-e), and it frames students as "agents to act, not simply objects to be acted upon" (BYU-Idaho, n.d.-c). This section measures current policy against the framework above and names the gaps to close.

What BYU-Idaho already does well.

- **It does not rely on AI detectors.** The academic-integrity policy is built on authorization, disclosure, and citation, not detection (BYU-Idaho, n.d.-b). This matches the strongest finding in the evidence (Section 2.2) and puts BYU-Idaho ahead of institutions still chasing detection.
- **It sets a clear default and requires disclosure.** AI may be used only when the instructor authorizes it, and use must be disclosed and cited (BYU-Idaho, n.d.-b). That is the spirit of stating permitted AI use per task.
- **It provides governed tools.** ChatGPT Edu, Microsoft Copilot, and Google Gemini run through institutional accounts with data-classification limits (BYU-Idaho, n.d.-d).
- **Its aims already echo this report.** Workforce readiness, "genuine understanding," and protecting the human and relational side of learning map onto the report's three aims (BYU-Idaho, n.d.-b, n.d.-c).

Gaps to close (measured against the framework).

7. **No secured-versus-open structure.** Policy says the instructor decides, but gives faculty no method for choosing which outcomes to secure AI-free and which to assess AI-augmented (Section 4). Decisions vary course to course.
8. **No shared AI-use scale.** Guidance is effectively binary (authorized or not) plus disclosure. There is no common, task-level scale, so "authorized" means different things in different courses (Section 2.2; recommendation 5).
9. **The assessment-design guidance is thin and hard to find.** The public integrity page tells faculty to "clearly define AI tool usage" and ensure "genuine understanding," but the how-to resources ("AI in the Syllabus," "Grading with AI") sit in a login-gated resource library, not on the public site, and are not visibly grounded in evidence.
10. **Everything is unit-level, not program-level.** Guidance stops at "check with your instructor." There is no program-level assessment mapping or a few secured key moments across a degree (Sections 2.6, 4.4).
11. **Disclosure is carrying too much weight.** Disclosure and citation are good, but declarations alone can fall flat (Section 2.2). The durable fix is assessment design, especially secured tasks, which the public policy underweights.

None of these require undoing what BYU-Idaho has built. They extend it: add a shared method and scale, surface the how-to, and lift assurance to the program level. Concrete website changes are in Appendix G.

8. Recommendations & Conclusion

Short-term wins (this term):

12. Stop using detection scores as evidence of misconduct; take an educative stance.

13. Have every instructor add one secured checkpoint (in-person or oral) that verifies individual understanding of core outcomes.
14. State permitted AI use for each task; allow AI freely for learning and practice.
15. Move some grade weight from the product to the process.
16. Adopt one common AI-use scale across the institution, such as the five-level AI Assessment Scale, so the per-task statements in step 3 use the same language everywhere (M. Perkins et al., 2024).

Long-term adjustments (1 to 2 years):

17. Move assessment security to the program level: map outcomes, secure key moments, replace single high-stakes exams with multiple points of evidence.
18. Fund team-based redesign and a sustained faculty resource repository with clear governance.
19. Run a program-level pilot, evaluating validity, workload, student experience, and teacher-student relationships.

For BYU-Idaho specifically, update the AI website (byui.edu/ai) so faculty can act on this framework: publish the secured-versus-open method and a shared AI-use scale, surface the login-gated how-to resources, and add program-level guidance. Concrete site changes are in Appendix G.

These steps serve all three aims: they restore the validity of what a degree certifies; they use oral and interactive formats that deepen teacher-student relationships; and they ready graduates to use AI well in work, community, and home. The direction is clear and well-evidenced: stop trying to detect the unmeasurable; design assessment that measures learning we can trust. We ask the committee and administration to endorse the framework, resource the repository and team redesign, and authorize the pilot.

References

- Ardito, C. G. (2025). Generative AI detection in higher education assessments. *New Directions for Teaching and Learning*, 2025(182), 11-28. <https://doi.org/10.1002/tl.20624>
- Association of American Colleges and Universities. (n.d.). *VALUE rubrics*. AAC&U. <https://www.aacu.org/initiatives/value-initiative/value-rubrics>
- Baytas, C., & Ruediger, D. (2025). *Making AI generative for higher education: Adoption and challenges among instructors and researchers*. Ithaka S+R. <https://doi.org/10.18665/sr.322677>
- Bowen, J. A., & Watson, C. E. (2024). *Teaching with AI: A practical guide to a new era of human learning*. Johns Hopkins University Press.
- Brattli, H., Utne, A., & Lynch, M. (2026). Assessment validity in the age of generative AI: A natural experiment. *Informatics*, 13(4), Article 56. <https://doi.org/10.3390/informatics13040056>

BYU-Idaho. (n.d.-a). *About artificial intelligence*. <https://www.byui.edu/ai/about>

BYU-Idaho. (n.d.-b). *AI and academic integrity*. <https://www.byui.edu/ai/academics/integrity>

BYU-Idaho. (n.d.-c). *AI priorities*. <https://www.byui.edu/ai/priorities>

BYU-Idaho. (n.d.-d). *BYU-Idaho authorized AI tools*. <https://www.byui.edu/ai/tools>

BYU-Idaho. (n.d.-e). *Prophetic guidance*. <https://www.byui.edu/ai/prophetic-guidance>

Carnegie Mellon University, Eberly Center. (2023). *The GAITAR initiative*. <https://www.cmu.edu/teaching/gaitar/index.html>

Church, I. M., Drake, L., & Harris, M. (2025). *Using LLMs to support assessment of student work in higher education: A viva voce simulator* [Preprint]. arXiv. <https://arxiv.org/abs/2511.05530>

Corbin, T., Bearman, M., Boud, D., & Dawson, P. (2026). The wicked problem of AI and assessment. *Assessment & Evaluation in Higher Education*, 51(4), 736-752. <https://doi.org/10.1080/02602938.2025.2553340>

Corbin, T., Bearman, M., Boud, D., Crawford, N., Dawson, P., Fawns, T., Henderson, M., Lodge, J., Luo, J., Matthews, K., Nicola-Richmond, K., Nieminen, J. H., Pepperell, N., Swiecki, Z., Tai, J., & Walton, J. (2025). Assessment after artificial intelligence: The research we should be doing. *Journal of University Teaching and Learning Practice*, 22(7), Article 2533. <https://doi.org/10.53761/w3x5y804>

Council of Regional Accrediting Commissions. (2025). *Statement on the use of artificial intelligence (AI) to advance learning, evaluation, and recognition*. C-RAC. <https://www.c-rac.org/post/c-rac-statement-on-the-use-of-artificial-intelligence-ai-to-advance-learning-evaluation-and-recognition>

Curtis, G. J. (2025). The two-lane road to hell is paved with good intentions. *Higher Education Research & Development*, 44(8), 2151-2158. <https://doi.org/10.1080/07294360.2025.2476516>

Deep, P. D., Edgington, W. D., Ghosh, N., & Rahaman, M. S. (2025). Evaluating the effectiveness and ethical implications of AI detection tools in higher education. *Information*, 16(10), Article 905. <https://doi.org/10.3390/info16100905>

Desai, H. (2025). *What's worth measuring? The future of assessment in the AI age*. UNESCO IdeasLAB. <https://www.unesco.org/en/articles/whats-worth-measuring-future-assessment-ai-age>

EDUCAUSE. (2025). *2025 EDUCAUSE AI landscape study*. <https://library.educause.edu/resources/2025/2/2025-educause-ai-landscape-study>

Freeman, J. (2025). *Student generative AI survey 2025* (HEPI Policy Note 61). Higher Education Policy Institute. <https://www.hepi.ac.uk/reports/student-generative-ai-survey-2025/>

Furze, L., Perkins, M., Roe, J., & MacVaugh, J. (2024). The AI Assessment Scale (AIAS) in action: A pilot implementation of GenAI-supported assessment. *Australasian Journal of Educational Technology*, 40(4), 38-55. <https://doi.org/10.14742/ajet.9434>

- Hung, J.-T., Cui, C., Popescu, D. M., Chatterjee, S., & Starner, T. (2024). Socratic Mind: Scalable oral assessment powered by AI. In *Proceedings of the Eleventh ACM Conference on Learning @ Scale*. Association for Computing Machinery. <https://doi.org/10.1145/3657604.3664661>
- Kofinas, A. K., Tsay, C. H.-H., & Pike, D. (2025). The impact of generative AI on academic integrity of authentic assessments within a higher education context. *British Journal of Educational Technology*, 56(6), 2522-2549. <https://doi.org/10.1111/bjet.13585>
- Laidlaw, P. (2026). Supporting assessment redesign in the age of AI: A critical practice-based framework for equitable change. *Teaching in Higher Education*. Advance online publication. <https://doi.org/10.1080/13562517.2026.2670362>
- Liang, W., Yuksekgonul, M., Mao, Y., Wu, E., & Zou, J. (2023). GPT detectors are biased against non-native English writers. *Patterns*, 4(7), Article 100779. <https://doi.org/10.1016/j.patter.2023.100779>
- Lodge, J. M., Howard, S., Bearman, M., & Dawson, P. (2023). *Assessment reform for the age of artificial intelligence*. Tertiary Education Quality and Standards Agency. <https://www.teqsa.gov.au/guides-resources/resources/corporate-publications/assessment-reform-age-artificial-intelligence>
- Martin, A. F., Tubaltseva, S., Harrison, A., & Rubin, G. J. (2025). Participatory co-design and evaluation of a novel approach to generative AI-integrated coursework assessment in higher education. *Behavioral Sciences*, 15(6), Article 808. <https://doi.org/10.3390/bs15060808>
- Nicola-Richmond, K., Dawson, P., Partridge, H., & Macfarlane, S. (2025). It takes a village: Program-wide approaches to redesigning assessment in a time of generative artificial intelligence (GenAI). *Journal of University Teaching and Learning Practice*, 22(7), Article 1711. <https://doi.org/10.53761/zpp2ja61>
- Pallant, J. L., Blijlevens, J., Campbell, A., & Jopp, R. (2025). Mastering knowledge: The impact of generative AI on student learning outcomes. *Studies in Higher Education*, 51(4), 714-735. <https://doi.org/10.1080/03075079.2025.2487570>
- Perkins, G. (2026). Resilient assessment in the age of AI: Authentic design and the case for verbal examinations in business education. *Assessment & Evaluation in Higher Education*. Advance online publication. <https://doi.org/10.1080/02602938.2026.2644516>
- Perkins, M., Furze, L., Roe, J., & MacVaugh, J. (2024). The Artificial Intelligence Assessment Scale (AIAS): A framework for ethical integration of generative AI in educational assessment. *Journal of University Teaching and Learning Practice*, 21(6). <https://doi.org/10.53761/q3azde36>
- Slimi, Z. (2026). A systematic critical review of generative AI's impact on authorship, pedagogy, and integrity (2023 to 2025). *Frontiers in Education*, 11, Article 1769680. <https://doi.org/10.3389/feduc.2026.1769680>
- Stokkink, P. (2025). *The impact of AI on educational assessment: A framework for constructive alignment* [Preprint]. arXiv. <https://arxiv.org/abs/2506.23815>

Tertiary Education Quality and Standards Agency. (2025). *Enacting assessment reform in a time of artificial intelligence*. TEQSA. <https://www.teqsa.gov.au/guides-resources/resources/corporate-publications/enacting-assessment-reform-time-artificial-intelligence>

University of London. (2025). *Ensuring assessment survives AI: A menu of tools and resources*. <https://onlinelearning.london.ac.uk/2025/07/01/menu-of-tools-and-resources/>

University of Massachusetts Amherst, Center for Teaching and Learning. (n.d.). *How do I redesign assignments and assessments for an AI-impacted world?* <https://www.umass.edu/ctl/how-do-i-redesign-assignments-and-assessments-ai-impacted-world>

University of Melbourne. (n.d.). *Interactive oral assessments*. Centre for the Study of Higher Education. <https://melbourne-cshe.unimelb.edu.au/resources/topics/assessment-and-feedback/specific-help/interactive-oral-assessments>

University of Pittsburgh, University Center for Teaching and Learning. (2023). *Guidance on AI-detection software*. University Times. <https://www.utimes.pitt.edu/news/teaching-center-doesn-t>

University of Sydney. (2024). *The Sydney assessment framework: A two-lane approach*. Teaching@Sydney. <https://educational-innovation.sydney.edu.au/teaching@sydney/the-sydney-assessment-framework/>

University of Sydney. (2025). *What we learned from 400+ interactive oral assessments in STEM*. Teaching@Sydney. <https://educational-innovation.sydney.edu.au/teaching@sydney/what-we-learned-from-400-interactive-oral-assessments-in-stem/>

Urquhart, L., & Ngo, X. M. (2026). Changing EAP assessment practices in the age of generative artificial intelligence: The case of Scottish higher education institutions. *Journal of English for Academic Purposes*, 79, Article 101609. <https://doi.org/10.1016/j.jeap.2025.101609>

Vanderbilt University. (2023). *Guidance on AI detection and why we're disabling Turnitin's AI detector*. Brightspace. <https://www.vanderbilt.edu/brightspace/2023/08/16/guidance-on-ai-detection-and-why-were-disabling-turnitins-ai-detector/>

Weber-Wulff, D., Anohina-Naumeca, A., Bjelobaba, S., Foltýnek, T., Guerrero-Dib, J., Popoola, O., Šigut, P., & Waddington, L. (2023). Testing of detection tools for AI-generated text. *International Journal for Educational Integrity*, 19(1), Article 26. <https://doi.org/10.1007/s40979-023-00146-z>

Weng, X., Xia, Q., Gu, M., Rajaram, K., & Chiu, T. K. F. (2024). Assessment and learning outcomes for generative AI in higher education: A scoping review on current research status and trends. *Australasian Journal of Educational Technology*, 40(6), 37-55. <https://doi.org/10.14742/ajet.9540>

Appendix A. Discipline Worked Examples (protect-vs-augment maps)

Cluster	Secure AI-free (verify directly)	Open / AI-augmented (assess judgment)
STEM	In-person problem-solving; oral defense of derivations	AI-assisted modeling or coding with written critique of AI output
Humanities	Seminar viva; in-class close-reading and writing	AI-assisted drafting with annotated reflection on edits and choices
Professional/Clinical	OSCE-style interactive orals; simulated skill checks	AI-supported care or plan drafting with justification and risk check
Social Sciences	Oral defense of method and analysis	AI-assisted data work with validity and ethics commentary

Appendix B. Assignment & Rubric Templates (patterns)

20. **Secured oral checkpoint:** 10 to 15 minutes; the student defends a piece of their own open work; rubric weights reasoning, accuracy, and response to probing questions.
21. **Process-weighted artifact:** student submits drafts, prompts, and a decision log; at least 30% of the grade rests on process evidence.
22. **AI-critique task:** student improves and critiques an AI-generated draft; graded on judgment, not polish.

Appendix C. Directory of Leading Figures & Organizations

People: Phillip Dawson, Margaret Bearman, David Boud, Rola Ajjawi, Joanna Tai (CRADLE, Deakin University); Jason Lodge (University of Queensland); Sarah Howard (University of Wollongong); Thomas Corbin (Deakin); Danny Liu (University of Sydney); Maha Bali (American University in Cairo). **US people:** José Antonio Bowen and C. Edward Watson (authors of Teaching with AI; Watson at AAC&U). **US teaching centers:** University of Michigan (CRLT), Vanderbilt (Center for Teaching), Carnegie Mellon (Eberly Center), Stanford (Teaching Commons). **Organizations:** CRADLE; TEQSA; QAA; Jisc; Advance HE; EDUCAUSE; AAC&U; Russell Group.

Appendix D. Training/Redesign Evaluation Instrument (outline)

Kirkpatrick-style: (1) faculty reaction; (2) learning, can map outcomes to secured or open tasks; (3) application, redesigned at least one assessment; (4) results: validity, workload, student experience, and relationship measures.

Appendix E. Sample Course AI-Policy Language (starter)

"This course distinguishes *secured* tasks, completed without AI to verify your individual understanding, from *open* tasks, where you may use AI and will be assessed on your judgment, process, and critique of AI output. The permitted use of AI is stated for each assignment. You may use AI freely to study and practice. AI-detection scores will not be used as evidence of misconduct."

Each task states an allowed AI level on a five-level scale, adapted from the AI Assessment Scale (M. Perkins et al., 2024): (1) No AI, work is fully your own; (2) AI Planning, AI for brainstorming and outlining only; (3) AI Collaboration, AI used and disclosed within stated limits; (4) Full AI, AI used freely with attribution; (5) AI Exploration, AI used and then critically evaluated as part of the task.

Appendix F. Emerging AI Tools for Assessment (named examples, not endorsements)

Evidence for these is early and often vendor-claimed. Treat them as tools to pilot, not proven measures of learning.

- **Socratic Mind (Georgia Tech, with UC San Diego).** AI agent that runs Socratic-style oral assessments, asking follow-up questions; a 2024 pilot plus a 2025 developer-reported impact preprint.
- **InStage.** AI voice agent for rehearsing presentations and verbal scenarios; vendor-reported outcomes only.
- **Microsoft Copilot and general LLMs.** Used as the mandated draft-and-critique tool in a co-designed, AI-integrated assessment (Martin et al., 2025).
- **University of Sydney practice bot.** Voice-enabled bot that lets students rehearse before interactive orals.
- **AI viva-voce simulators (research prototypes).** Proof-of-concept systems that read a submission and ask four to five oral questions; not yet evaluated deployments.

Appendix G. Recommendations to Improve the AI Website (byui.edu/ai)

Concrete changes to align the site with this report, easiest first. Each maps to a recommendation or section.

23. **Publish a faculty "Assessment in the age of AI" pathway.** A public page that walks faculty through the secured-versus-open decision (Section 4.3), with examples. Today that how-to lives behind login.
24. **Post one shared AI-use scale.** Put a single five-level scale (Appendix E) on the public syllabus guidance so "authorized" means the same thing across courses.
25. **Surface the gated resources.** Move or mirror "AI in the Syllabus" and "Grading with AI" from the login-gated resource library onto the public /ai/academics area, and ground them in the evidence.
26. **Add program-level guidance.** A short page on mapping outcomes and securing a few key moments across a program, not only per course (Sections 2.6, 4.4).
27. **State the detection stance explicitly.** Say plainly that AI-detection scores are not used as evidence of misconduct, so no course adopts detectors ad hoc. BYU-Idaho already omits detection; make it explicit.
28. **Add ready-to-use templates.** Sample AI-use statements at each scale level, a secured oral protocol, and a process-weighted rubric (Appendices B, D, E).
29. **Reorganize navigation for faculty.** Separate three clear paths: Teaching and Assessment, Tools, and Data and Governance, with a "start here for course design" entry. Current navigation mixes governance and pedagogy.
30. **Connect the faith-centered framing to design.** Show that protecting durable skills means securing certain outcomes AI-free, and that oral and interactive assessment strengthens the teacher-student relationships the guidance already values.